

Deep Slither.io

A Deep Q Learning Approach to Slither.io

Brian Voter [bv85]

The Game



- Slither.io – a massively multiplayer browser game
 - Mouse controls a worm
 - Consuming glowing “food” pellets increases worm size and score
 - Cut off other worms to make them disintegrate into food
 - Goal: Become as large as possible, as fast possible
-

Objective and Constraints

- Learn to play optimally using only the game pixels as inputs
 - Reward signal is calculated each transition: $r = \text{score}(s') - \text{score}(s)$
 - Inspired by Minh et al. 2013: Playing Atari with Deep Reinforcement Learning
 - Key Differences:
 - Not an offline emulator – can't control flow of states
 - Instead must poll current state from the browser at a regular interval
-

High-level Approach

- Game is controlled using browser automation software
 - Neural Net accepts pixel matrix as input
 - Output: predicted quality of each action for the state
 - The action of the highest quality is performed, yielding a reward and transitioning to a new state
-

Background: Q-Learning

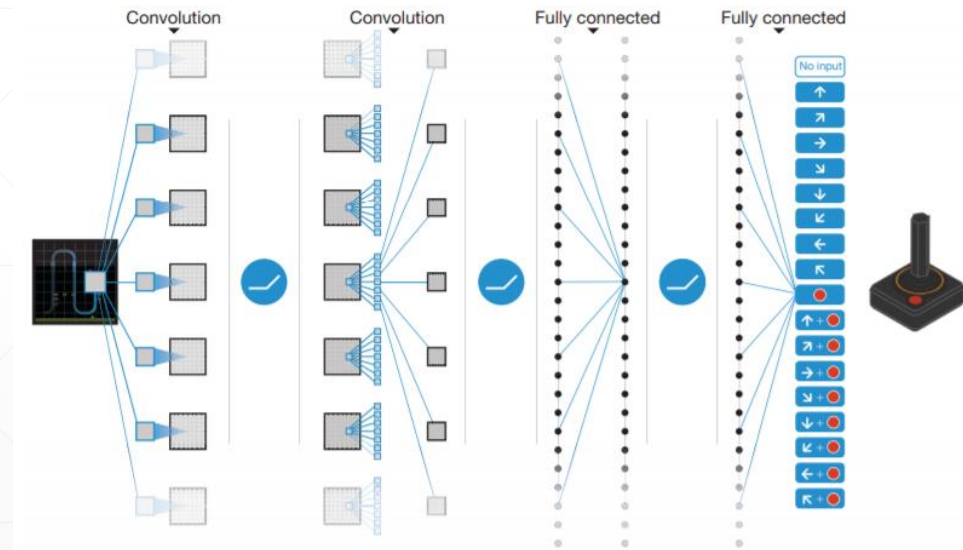
- Foundation of approach is Q-learning
 - Basic Q function: $Q(s, a)$ is the “quality” of performing action a in state s
 - **Problem:** Slither.io has infinitely many states
 - Can't learn quality of all possible states
 - Need to approximate
-

Background: DQN

- **Solution:** Approximate $Q(s, a) \approx Q(s, a; \theta)$
 - Right side uses neural network with parameters θ to approximate left side
 - Introduced by Minh et al. 2013
 - Deep Q Network = DQN
 - DQN is the foundation for my approach
-

Background: DQN Architecture

- Original DQN accepts $[4 \times 84 \times 84]$ images
- Faster processing because images are downscaled from $[210 \times 160]$
- 4 images are stacked on the channel-dimension to give state history
 - Consider single *Pong* frame
 - For Atari 2600 games, changes POMDP \rightarrow MDP



Mnih et al. 2015, *Nature*

Background: Experience Replay

- **Problem:** as a DQN is trained on fresh states, it will “forget” transitions it learned earlier
 - **Solution:** *Experience replay*
 - Don't just train on fresh states
 - With a replay buffer B holding up to N states, instead train on sampled batch: $b \sim B$
 - Default mode is uniform random sampling
-

Background: Prioritized Experience Replay

- Replay buffer B holding up to N states, train on sampled batch: $b \sim B$
- The probability that a transition is sampled depends on its “surprise factor” – how much the network learned from the transition
 - $\text{error} = Q(s, a) - \left(r + \gamma \max_{a'} Q(s', a') \right)$
 - $\text{priority} = |\text{error}| + \epsilon$
- But more difficult to store states and sample efficiently:
 - Replacing e.g. a ring buffer, authors recommend a *sum tree* achieving $O(\log n)$ sampling
- Used in Slither system

Background: Double DQN

- **Problem:** DQN can significantly overestimate Q values
- **Solution:** Double DQN
 - Use a target network, initialize parameters: $\theta^- = \theta$ at startup
 - At (long) regular intervals, copy $\theta^- \leftarrow \theta$ (to the target network from the online network)
 - Replace target values: $y_t^{\text{DQN}} \equiv R_{t+1} + \gamma \max_a Q(S_{t+1}, a; \theta_t)$
with: $y_t^{\text{DoubleQ}} \equiv R_{t+1} + \gamma Q(S_{t+1}, \operatorname{argmax}_a Q(S_{t+1}, a; \theta_t); \theta_t^-)$.
 - That is, *choose* the best action for state S_{t+1} using online network, but *evaluate* it using the target network
 - Reduces overestimation and is implemented in Slither system

Slither System Overview

- Written in Java
 - Browser control with Selenium WebDriver
 - Neural network implemented in DL4J / ND4J
 - Image capture / processing with OpenCV/JavaCV
 - Trained with 16GB RAM system with CUDA on NVIDIA 980TI GPU
-

Neural Network Architecture

- Similar to that of the original DQN
 - Larger networks take longer to train
 - Input: [1x120x60]
 - Convolutional Layers
 - Layer 1: 16 8×8 filters, stride 4
 - Layer 2: 32 4×4 filters, stride 2
 - Layer 3: Fully connected, 256 outputs
 - Output: numOutputs = numActions (explained later)
-

Neural Network Architecture

- Why just one input image?
 - Using multiple images for each transition seems to require either more memory (preprocess once) or recurring computation (combine images at training time)
 - Neither is desirable
 - Direction of a worm = direction of head
 - Why the different input size?
 - DQN used 84x84, apparently because their NN library preferred square images
 - Instead we have a similar image area and retain the same aspect ratio
-

Neural Network Details

- Output layer has linear activation
 - All other layers have ReLU (rectified linearity) activation
 - Stochastic Gradient Descent
 - Updater: Adam ($\alpha = 0.0005$, $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 1 \times 10^{-8}$)
 - As recommended by Kingma and Ba (2014) except for α
 - Loss: Mean Squared Error
 - Good alternatives for further study include Huber Loss and LogCosh Loss
-

Neural Network Details

- Weights are initialized according to He et al. (2015):
 - Issue: Xavier initializer assumes activations are linear (they aren't)
 - Can slow or prevent convergence
 - Proposed ReLU initializer \sim Normal with mean 0, stdev $\sqrt{\frac{2}{\hat{n}_l}}$ where $\hat{n}_l = k_l^2 d_l$, l is the layer index, k is the filter size, and d is the number of filters
-

Image Processing

- Screenshots of the game are taken rapidly using JavaCV's FFmpeg extension, as opposed to Java's Robot class
 - Robot class too slow for realtime inference
 - Images are downscaled to $[120 \times 60]$
 - Raw images roughly $[1800 \times 900]$
 - 3 Channels are reduced to 1
 - Pixel values are normalized from $[0, 255]$ to $[0, 1]$
-

Action Set Discretization

- **Problem:** Slither.io admits a nearly continuous set of inputs
 - $1800 \times 900 \times 2 = 3,240,000$ actions.
 - Recall `numOutputs = numActions`
-

Action Set Discretization

“Framing” Approach



- Preselect a frame of absolute coordinates
- Mouse “jumps” to coordinates
- Leads to jittery movements

“Sliding” Approach



- $a_x, a_y \in \{-100, 0, +100\}, a_{\text{boost}} \in \{0,1\}$
- Mouse slides ≤ 100 pixels in any direction per timestep
- Encourages more commitment to one direction

Gameplay Algorithm

- See pseudocode in report
 - Agent acts under ϵ - greedy policy
 - ϵ is linearly annealed from 1.0 to ~ 0.1 over thousands of timesteps
 - If $\text{rand}() < \epsilon$: use random action; else: use best action
 - Every ~ 5 timesteps: train on sample transition batch
 - Popular batch size = 32
 - Every ~ 1000 timesteps: copy $\theta_t^- \leftarrow \theta_t$
-

Evaluation

- Unfortunately, agent performs below human skill level
 - Agent scores generally < 1000 points
 - A small convenience sample of human scores was obtained by joining multiple game sessions
 - Observed score of players at position 1 of the leaderboard
 - $N = 7$
 - Mean = 34,157
 - Median = 32,099
 - Stdev: 16,640
 - However YouTube videos allege scores of 100,000+
-

Conclusion

- Agent performance was worse than one might hope
 - Agent has not discovered the crucial “cut-off” and “encircling” strategies
 - I discovered that Mnih et al. trained their Atari 2600 agent, for each game, on 10M frames
 - Slither system can process ~ 12,500 frames / hour, requiring 30+ days to reach 10M
 - However, learning efficiency (at least on Atari 2600) can be further improved as shown by Hessel et al. (2017) in the *Rainbow* agent
 - Using all of: double learning, prioritized replay, dueling networks, multi-step learning, distributional RL, and noisy nets
 - Independent DQN extensions can be combined for greater overall performance
-

Conclusion

- A unique test-bed for RL algorithms was created
 - The Slither system is distinct from the Atari literature as the game has highly limited observability
 - Agent only sees a small portion of the map at a time
 - Further study: is DQN or a variant thereof well suited for partially observable games?
 - How distinct is flickering Pong from Slither.io?
 - See also: recurrent DQN (Hausknecht et al. 2015)
-

Complete References

- [1] V. Mnih, K. Kavukcuoglu, D. Silver, A. Graves, I. Antonoglou, D. Wierstra, M. Riedmiller, Playing Atari with Deep Reinforcement Learning, ArXiv13125602 Cs. (2013). <http://arxiv.org/abs/1312.5602> (accessed May 19, 2018).
 - [2] T. Schaul, J. Quan, I. Antonoglou, D. Silver, Prioritized Experience Replay, ArXiv151105952 Cs. (2015). <http://arxiv.org/abs/1511.05952> (accessed May 19, 2018).
 - [3] S. Zhang, R.S. Sutton, A Deeper Look at Experience Replay, ArXiv171201275 Cs. (2017). <http://arxiv.org/abs/1712.01275> (accessed May 19, 2018).
 - [4] H. van Hasselt, A. Guez, D. Silver, Deep Reinforcement Learning with Double Q-learning, ArXiv150906461 Cs. (2015). <http://arxiv.org/abs/1509.06461> (accessed May 19, 2018).
 - [5] D.P. Kingma, J. Ba, Adam: A Method for Stochastic Optimization, ArXiv14126980 Cs. (2014). <http://arxiv.org/abs/1412.6980> (accessed May 20, 2018).
 - [6] K. He, X. Zhang, S. Ren, J. Sun, Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification, ArXiv150201852 Cs. (2015). <http://arxiv.org/abs/1502.01852> (accessed May 20, 2018).
 - [7] M. Hausknecht, P. Stone, Deep Recurrent Q-Learning for Partially Observable MDPs, ArXiv150706527 Cs. (2015). <http://arxiv.org/abs/1507.06527> (accessed May 20, 2018).
 - [8] M. Hessel, J. Modayil, H. van Hasselt, T. Schaul, G. Ostrovski, W. Dabney, D. Horgan, B. Piot, M. Azar, D. Silver, Rainbow: Combining Improvements in Deep Reinforcement Learning, ArXiv171002298 Cs. (2017). <http://arxiv.org/abs/1710.02298> (accessed May 20, 2018).
-